
ARQUITECTURA DE COMPUTADORES

Tema 2: Instrucciones: El Lenguaje del Computador

Mary Jane Irwin (www.cse.psu.edu/~mji)

[Adapted from *Computer Organization and Design, 4th Edition*,
Patterson & Hennessy, © 2008, MK]

Review : Evaluating ISAs

❑ Design-time metrics:

- | Can it be implemented? With what performance, at what costs (design, fabrication, test, packaging), with what power, with what reliability?
- | Can it be programmed? Ease of compilation?

❑ Static Metrics:


- | How many bytes does the program occupy in memory?

❑ Dynamic Metrics:

- | How many instructions are executed? How many bytes does the processor fetch to execute the program?
- | How many clocks are required per instruction?
- | How "lean" a clock is practical?

Best Metric: Time to execute the program!

depends on the instructions set, the processor organization, and compilation techniques.

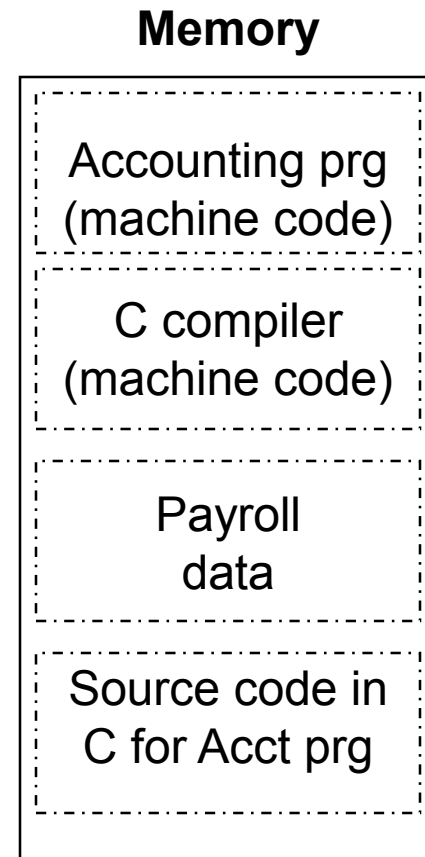


Two Key Principles of Machine Design

1. Instructions are represented as numbers and, as such, are indistinguishable from data
2. Programs are stored in alterable memory (that can be read or written to) just like data

❑ Stored-program concept

- | Programs can be shipped as files of binary numbers – **binary compatibility**
- | Computers can inherit ready-made software provided they are compatible with an existing ISA – leads industry to align around a small number of ISAs




MIPS-32 ISA

❑ Instruction Categories

- | Computational
- | Load/Store
- | Jump and Branch
- | Floating Point
 - coprocessor
- | Memory Management
- | Special

Registers



3 Instruction Formats: **all 32 bits wide**

op	rs	rt	rd	sa	funct	R format
op	rs	rt	immediate			I format
op	jump target					J format

MIPS (RISC) Design Principles

❑ Simplicity favors regularity

- | fixed size instructions
- | small number of instruction formats
- | opcode always the first 6 bits

❑ Smaller is faster

- | limited instruction set
- | limited number of registers in register file
- | limited number of addressing modes

❑ Make the common case fast

- | arithmetic operands from the register file (load-store machine)
- | allow instructions to contain immediate operands

❑ Good design demands good compromises

- | three instruction formats

MIPS Arithmetic Instructions

- ❑ MIPS assembly language arithmetic statement

add \$t0, \$s1, \$s2

sub \$t0, \$s1, \$s2

- ❑ Each arithmetic instruction performs **one** operation
- ❑ Each specifies exactly **three** operands that are all contained in the datapath's register file (\$t0, \$s1, \$s2)

destination ← source1 **op** source2

- ❑ Instruction Format (**R** format)



MIPS Instruction Fields

- ❑ MIPS fields are given names to make them easier to refer to

op	rs	rt	rd	shamt	funct
----	----	----	----	-------	-------

op	6-bits	opcode that specifies the operation
rs	5-bits	register file address of the first source operand
rt	5-bits	register file address of the second source operand
rd	5-bits	register file address of the result's destination
shamt	5-bits	shift amount (for shift instructions)
funct	6-bits	function code augmenting the opcode

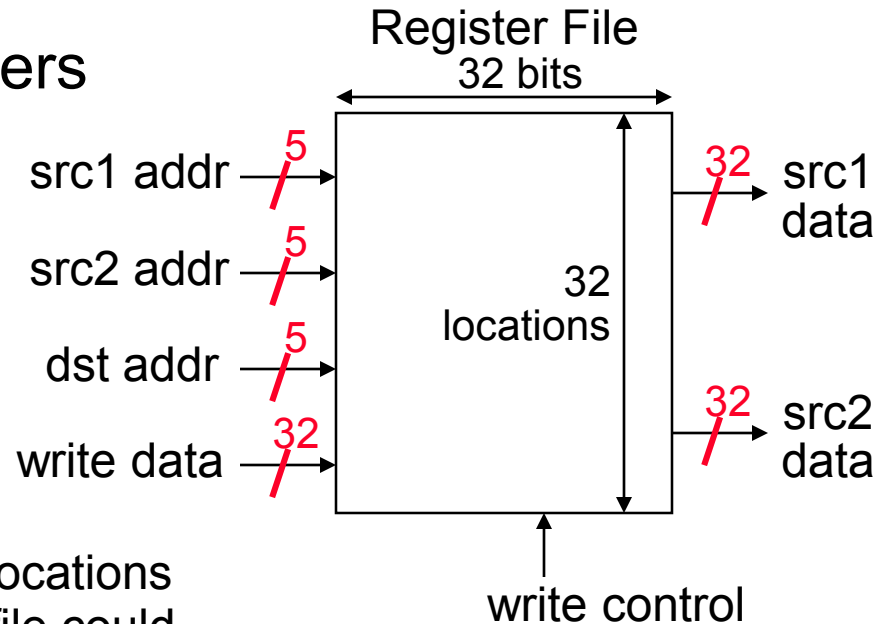
MIPS Register File

❑ Holds thirty-two 32-bit registers

- | Two read ports and
- | One write port

❑ Registers are

- | **Faster** than main memory
 - But register files with more locations are slower (e.g., a 64 word file could be as much as 50% slower than a 32 word file)
 - Read/write port increase impacts speed quadratically
- | Easier for a compiler to use
 - e.g., $(A*B) - (C*D) - (E*F)$ can do multiplies in any order vs. stack
- | Can hold variables so that
 - code density improves (since register are named with fewer bits than a memory location)



Aside: MIPS Register Convention

Name	Register Number	Usage	Preserve on call?
\$zero	0	constant 0 (hardware)	n.a.
\$at	1	reserved for assembler	n.a.
\$v0 - \$v1	2-3	returned values	no
\$a0 - \$a3	4-7	arguments	yes
\$t0 - \$t7	8-15	temporaries	no
\$s0 - \$s7	16-23	saved values	yes
\$t8 - \$t9	24-25	temporaries	no
\$gp	28	global pointer	yes
\$sp	29	stack pointer	yes
\$fp	30	frame pointer	yes
\$ra	31	return addr (hardware)	yes

MIPS Memory Access Instructions

- ❑ MIPS has two basic **data transfer** instructions for accessing memory


`lw $t0, 4($s3) #load word from memory`

`sw $t0, 8($s3) #store word to memory`

- ❑ The data is loaded into (lw) or stored from (sw) a register in the register file – a 5 bit address
- ❑ The memory address – a 32 bit address – is formed by adding the contents of the **base address register** to the **offset** value
 - A 16-bit field meaning access is limited to memory locations within a region of $\pm 2^{13}$ or 8,192 words ($\pm 2^{15}$ or 32,768 bytes) of the address in the base register


Machine Language - Load Instruction

Load/Store Instruction Format (I format):




$$24_{10} + \$s3 =$$

$$\begin{array}{r}
 \dots 0001\ 1000 \\
 + \dots 1001\ 0100 \\
 \hline
 \dots 1010\ 1100 = \\
 0x120040ac
 \end{array}$$



Byte Addresses

- ❑ Since 8-bit bytes are so useful, most architectures address individual **bytes** in memory
 - **Alignment restriction** - the memory address of a **word** must be on natural word boundaries (a multiple of 4 in MIPS-32)
- ❑ **Big Endian:** leftmost byte is word address
IBM 360/370, Motorola 68k, **MIPS**, Sparc, HP PA
- ❑ **Little Endian:** rightmost byte is word address
Intel 80x86, DEC Vax, DEC Alpha (Windows NT)



Aside: Loading and Storing Bytes

- ❑ MIPS provides special instructions to move bytes

```
lb    $t0, 1($s3)    #load byte from memory
```

```
sb    $t0, 6($s3)    #store byte to memory
```

0x28	19	8	16 bit offset
------	----	---	---------------

- ❑ What 8 bits get loaded and stored?

- load byte places the byte from memory in the rightmost 8 bits of the destination register
 - what happens to the other bits in the register?
- store byte takes the byte from the rightmost 8 bits of a register and writes it to a byte in memory
 - what happens to the other bits in the memory word?

MIPS Immediate Instructions

- ❑ Small constants are used often in typical code
- ❑ Possible approaches?
 - put “typical constants” in memory and load them
 - create hard-wired registers (like \$zero) for constants like 1
 - have special instructions that contain constants !

`addi $sp, $sp, 4` $\# \$sp = \$sp + 4$

`slti $t0, $s2, 15` $\# \$t0 = 1 \text{ if } \$s2 < 15$

- ❑ Machine format (I format):

0x0A	18	8	0x0F
------	----	---	------

- ❑ The constant is kept **inside** the instruction itself!
 - Immediate format **limits** values to the range $+2^{15}-1$ to -2^{15}

Aside: How About Larger Constants?

- ❑ We'd also like to be able to load a 32 bit constant into a register, for this we must use two instructions
- ❑ a new "load upper immediate" instruction

```
lui $t0, 1010101010101010
```

16	0	8	1010101010101010 ₂
----	---	---	-------------------------------

- ❑ Then must get the lower order bits right, use

```
ori $t0, $t0, 1010101010101010
```


1010101010101010	0000000000000000
------------------	------------------

0000000000000000	1010101010101010
------------------	------------------

1010101010101010	1010101010101010
------------------	------------------

Review: Unsigned Binary Representation

Hex	Binary	Decimal
0x00000000	0...0000	0
0x00000001	0...0001	1
0x00000002	0...0010	2
0x00000003	0...0011	3
0x00000004	0...0100	4
0x00000005	0...0101	5
0x00000006	0...0110	6
0x00000007	0...0111	7
0x00000008	0...1000	8
0x00000009	0...1001	9
	...	
0xFFFFFFFFC	1...1100	$2^{32} - 4$
0xFFFFFFFDD	1...1101	$2^{32} - 3$
0xFFFFFFFEE	1...1110	$2^{32} - 2$
0xFFFFFFFFF	1...1111	$2^{32} - 1$



Review: Signed Binary Representation

$$-2^3 =$$

$$-(2^3 - 1) =$$

2'sc binary	decimal
1000	-8
1001	-7
1010	-6
1011	-5
1100	-4
1101	-3
1110	-2
1111	-1
0000	0
0001	1
0010	2
0011	3
0100	4
0101	5
0110	6
0111	7

complement all the bits

0101

1011

and add a 1

and add a 1

0110

1010

complement all the bits

$$2^3 - 1 =$$

MIPS Shift Operations

- ❑ Need operations to **pack** and **unpack** 8-bit characters into 32-bit words
- ❑ Shifts move all the bits in a word left or right

`sll $t2, $s0, 8` `#$t2 = $s0 << 8 bits`

`srl $t2, $s0, 8` `#$t2 = $s0 >> 8 bits`

- ❑ Instruction Format (**R** format)

0		16	10	8	0x00
---	--	----	----	---	------

- ❑ Such shifts are called **logical** because they fill with **zeros**
 - | Notice that a 5-bit shamt field is enough to shift a 32-bit value $2^5 - 1$ or **31 bit positions**

MIPS Logical Operations

- There are a number of **bit-wise** logical operations in the MIPS ISA

`and $t0, $t1, $t2 #$t0 = $t1 & $t2`

`or $t0, $t1, $t2 #$t0 = $t1 | $t2`

`nor $t0, $t1, $t2 #$t0 = not($t1 | $t2)`

- Instruction Format (**R** format)

0	9	10	8	0	0x24
---	---	----	---	---	------

`andi $t0, $t1, 0xFF00 #$t0 = $t1 & ff00`

`ori $t0, $t1, 0xFF00 #$t0 = $t1 | ff00`

- Instruction Format (**I** format)

0x0D	9	8	0xFF00
------	---	---	--------

MIPS Control Flow Instructions

❑ MIPS conditional branch instructions:

```
bne $s0, $s1, Lbl1 #go to Lbl1 if $s0≠$s1
beq $s0, $s1, Lbl1 #go to Lbl1 if $s0=$s1
```

| Ex: if (i==j) h = i + j;

```
        bne $s0, $s1, Lbl1
        add $s3, $s0, $s1
Lbl1:    ...
```

❑ Instruction Format (I format):

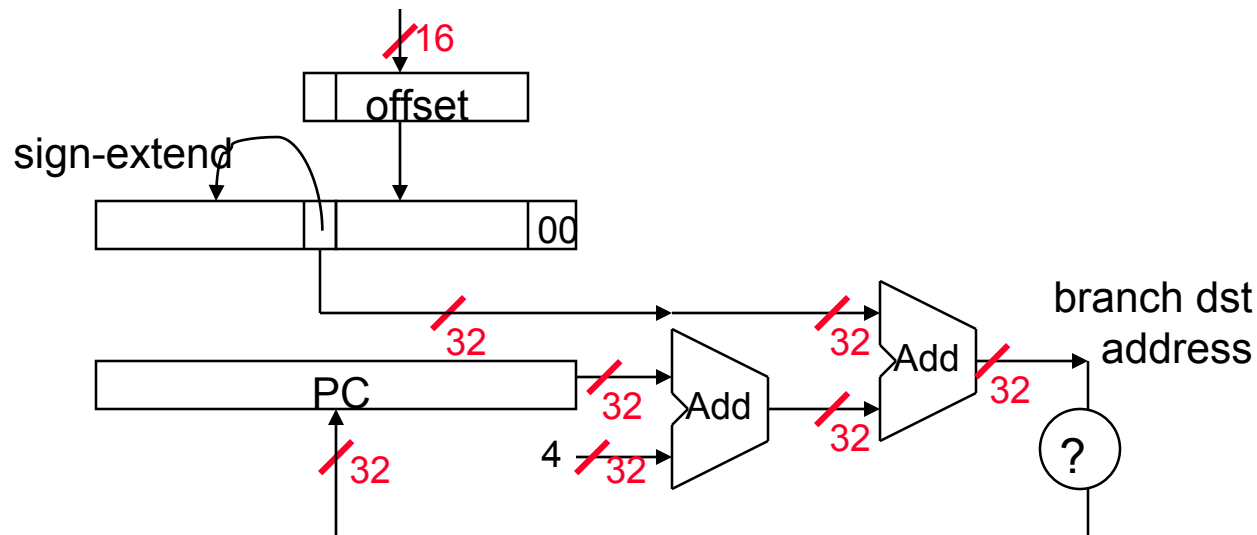
0x05	16	17	16 bit offset
------	----	----	---------------

❑ How is the branch destination address specified?

Specifying Branch Destinations

- ❑ Use a register (like in lw and sw) added to the 16-bit offset
 - | which register? Instruction Address Register (the PC)
 - its use is automatically **implied** by instruction
 - PC gets updated (PC+4) during the **fetch** cycle so that it holds the address of the next instruction
 - | limits the branch distance to **-2^{15} to $+2^{15}-1$** (word) instructions from the (instruction after the) branch instruction, but most branches are local anyway

from the low order 16 bits of the branch instruction



In Support of Branch Instructions

- ❑ We have `beq`, `bne`, but what about other kinds of branches (e.g., branch-if-less-than)? For this, we need yet another instruction, `slt`

- ❑ Set on less than instruction:

```
slt $t0, $s0, $s1    # if $s0 < $s1    then
                      # $t0 = 1          else
                      # $t0 = 0
```

- ❑ Instruction format (**R** format):

0	16	17	8		0x24
---	----	----	---	--	------

- ❑ Alternate versions of `slt`

```
slti $t0, $s0, 25    # if $s0 < 25 then $t0=1 ...
sltu $t0, $s0, $s1    # if $s0 < $s1 then $t0=1 ...
sltiu $t0, $s0, 25    # if $s0 < 25 then $t0=1 ...
```

Aside: More Branch Instructions

- ❑ Can use `slt`, `beq`, `bne`, and the fixed value of 0 in register `$zero` to **create** other conditions

- | less than `blt $s1, $s2, Label`

- `slt $at, $s1, $s2` `#$at set to 1 if`

- `bne $at, $zero, Label` `#$s1 < $s2`

- | less than or equal to `ble $s1, $s2, Label`

- | greater than `bgt $s1, $s2, Label`

- | great than or equal to `bge $s1, $s2, Label`

- ❑ Such branches are included in the instruction set as pseudo instructions - recognized (and expanded) by the assembler

- | Its why the assembler needs a reserved register (`$at`)

Bounds Check Shortcut

- ❑ Treating signed numbers as if they were unsigned gives a low cost way of checking if $0 \leq x < y$ (index out of bounds for arrays)

```
sltu $t0, $s1, $t2      # $t0 = 0 if
                        # $s1 > $t2 (max)
                        # or $s1 < 0 (min)
beq $t0, $zero, IOOB     # go to IOOB if
                        # $t0 = 0
```


- ❑ The key is that negative integers in two's complement look like large numbers in unsigned notation. Thus, an unsigned comparison of $x < y$ also checks if x is negative as well as if x is less than y .

Other Control Flow Instructions


- ❑ MIPS also has an unconditional branch instruction or **jump** instruction:

```
j    label        #go to label
```

- ## ❑ Instruction Format (J Format):



from the low order 26 bits of the jump instruction



Aside: Branching Far Away

- ❑ What if the branch destination is further away than can be captured in 16 bits?
- ❑ The assembler comes to the rescue – it inserts an unconditional jump to the branch target and inverts the condition

```
beq    $s0, $s1, L1
```

becomes


```
bne    $s0, $s1, L2
```

```
j      L1
```


```
L2 :
```

Addressing Modes Illustrated

1. Register addressing




2. Base (displacement) addressing




3. Immediate addressing



4. PC-relative addressing



5. Pseudo-direct addressing



MIPS Organization So Far

